

Report Project 3

Object Detection, Semantic Segmentation, and Instance Segmentation

I am using my 3 late days for assignment 3 (Project 2)

Name: *Jashanraj Singh Gosain*

Student Number: *301435386*

Late Days: *3 Used*

Part 1:

The following configuration was used in part1.

```
'''
# Set the configs for the detection part in here.
# TODO: approx 15 lines
'''

from detectron2.engine import DefaultTrainer

cfg = get_cfg()
cfg.OUTPUT_DIR = "{}output/".format(BASE_DIR)


cfg.merge_from_file(model_zoo.get_config_file("COCO-Detection/faster_rcnn_X_101_32x8d_FPN_3x.yaml"))
cfg.DATASETS.TRAIN = ("data_detection_train",)
cfg.DATASETS.TEST = ("data_detection_test",)
cfg.DATALOADER.NUM_WORKERS = 4
cfg.MODEL.WEIGHTS = model_zoo.get_checkpoint_url("COCO-Detection/faster_rcnn_X_101_32x8d_FPN_3x.yaml") # Let training initialize from model zoo
cfg.SOLVER.IMS_PER_BATCH = 2 # "batch size"
cfg.SOLVER.BASE_LR = 0.00025 # LR
cfg.SOLVER.MAX_ITER = 1500 # iterations
cfg.SOLVER.STEPS = [] # decay learning rate
cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE = 512 # "RoIHead batch size"
cfg.MODEL.ROI_HEADS.NUM_CLASSES = 1 # only has one class (Plane)
```

Modifications:

- Max Iterations: Increasing the number of iterations helped lower the loss. With 500 iterations, there was more scope of fine tuning to lower the loss further. Increasing iterations helped reduce the loss.
- Low learning rate has little scope of improvement. It converges to the final accuracy in fewer iterations. High learning rate starts with good accuracy but the accuracy fluctuates and loss starts to increase after convergence. Ideally 0.00025 shows steady improvement with little fluctuations.
- "faster_rcnn_X_101_32x8d_FPN_3x.yaml" is supposedly more accurate model to train.
- ***cfg.DATALOADER.NUM_WORKERS = 4***

It helps speed up the training by employing 4 process on separate GPU core in parallel.

total_loss
tag: total_loss



Final Accuracy: 68.668

```
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.366
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.569
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.420
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.263
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.453
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.687
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.017
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.141
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.400
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.263
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.493
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.781
```

[11/11 23:26:55 d2.evaluation.coco_evaluation]: Evaluation results for bbox:

AP	AP50	AP75	APs	APm	APl
36.606	56.865	41.969	26.251	45.258	68.668


OrderedDict([('bbox', {'AP': 36.6059251323455, 'AP50': 56.865410003200445, 'AP75': 41.96888714629425, 'APs': 26.250909023515234, 'APm': 45.25848723794878, 'APl': 68.66845198902169})])

Visualization of test samples

Sample 1



Sample 2



Sample 3



Ablation Study:

Increasing `cfg.MODEL.ROI_HEADS.BATCH_SIZE_PER_IMAGE` made the convergence faster:

- For 1024, it converged faster but it keeps fluctuating and loss doesn't reduce further.

Converge to total_loss of around 1.

```
[11/11 21:56:23 d2.engine.train_loop]: Starting training from iteration 0
[11/11 21:57:13 d2.utils.events]: eta: 0:49:40 iter: 19 total_loss: 4.372 loss_cls: 0.6757 loss_box_reg: 0.1734 loss_rpn_cls: 2.953 loss_rpn_loc: 0.5968 time: 2.3204 last_time: 1.7986 data_time:
[11/11 21:57:55 d2.utils.events]: eta: 0:46:14 iter: 39 total_loss: 2.861 loss_cls: 0.6328 loss_box_reg: 0.2277 loss_rpn_cls: 0.6847 loss_rpn_loc: 0.392 time: 2.2033 last_time: 1.7849 data_time:
[11/11 21:58:39 d2.utils.events]: eta: 0:45:57 iter: 59 total_loss: 1.453 loss_cls: 0.5575 loss_box_reg: 0.2259 loss_rpn_cls: 0.2968 loss_rpn_loc: 0.2953 time: 2.2052 last_time: 1.8540 data_time:
[11/11 21:59:17 d2.utils.events]: eta: 0:45:04 iter: 79 total_loss: 1.284 loss_cls: 0.4525 loss_box_reg: 0.1779 loss_rpn_cls: 0.2197 loss_rpn_loc: 0.3423 time: 2.1335 last_time: 1.8873 data_time:
[11/11 22:00:00 d2.utils.events]: eta: 0:45:19 iter: 99 total_loss: 1.178 loss_cls: 0.3733 loss_box_reg: 0.1898 loss_rpn_cls: 0.2865 loss_rpn_loc: 0.293 time: 2.1286 last_time: 1.5080 data_time:
[11/11 22:00:41 d2.utils.events]: eta: 0:45:39 iter: 119 total_loss: 0.9994 loss_cls: 0.326 loss_box_reg: 0.2075 loss_rpn_cls: 0.1971 loss_rpn_loc: 0.2059 time: 2.1221 last_time: 1.3430 data_time:
[11/11 22:01:21 d2.utils.events]: eta: 0:44:57 iter: 139 total_loss: 1.159 loss_cls: 0.3193 loss_box_reg: 0.2471 loss_rpn_cls: 0.2113 loss_rpn_loc: 0.3417 time: 2.1022 last_time: 1.3939 data_time:
[11/11 22:02:02 d2.utils.events]: eta: 0:44:12 iter: 159 total_loss: 1.107 loss_cls: 0.2695 loss_box_reg: 0.255 loss_rpn_cls: 0.2339 loss_rpn_loc: 0.269 time: 2.0940 last_time: 1.6294 data_time:
[11/11 22:02:49 d2.utils.events]: eta: 0:43:40 iter: 179 total_loss: 1.091 loss_cls: 0.2701 loss_box_reg: 0.2088 loss_rpn_cls: 0.206 loss_rpn_loc: 0.2952 time: 2.1218 last_time: 1.7894 data_time:
[11/11 22:03:33 d2.utils.events]: eta: 0:43:49 iter: 199 total_loss: 0.8985 loss_cls: 0.2118 loss_box_reg: 0.2224 loss_rpn_cls: 0.1865 loss_rpn_loc: 0.2643 time: 2.1307 last_time: 1.6979 data_time:
[11/11 22:04:12 d2.utils.events]: eta: 0:42:52 iter: 219 total_loss: 1.045 loss_cls: 0.2401 loss_box_reg: 0.1883 loss_rpn_cls: 0.1696 loss_rpn_loc: 0.2777 time: 2.1159 last_time: 1.6931 data_time:
[11/11 22:04:53 d2.utils.events]: eta: 0:42:13 iter: 239 total_loss: 1.043 loss_cls: 0.2789 loss_box_reg: 0.2342 loss_rpn_cls: 0.1817 loss_rpn_loc: 0.2386 time: 2.1099 last_time: 1.6793 data_time:
[11/11 22:05:36 d2.utils.events]: eta: 0:41:32 iter: 259 total_loss: 0.9977 loss_cls: 0.2717 loss_box_reg: 0.2624 loss_rpn_cls: 0.1773 loss_rpn_loc: 0.2364 time: 2.1102 last_time: 1.5884 data_time:
[11/11 22:06:17 d2.utils.events]: eta: 0:40:52 iter: 279 total_loss: 0.9007 loss_cls: 0.2625 loss_box_reg: 0.2995 loss_rpn_cls: 0.1686 loss_rpn_loc: 0.1743 time: 2.1063 last_time: 2.0552 data_time:
[11/11 22:07:05 d2.utils.events]: eta: 0:40:35 iter: 299 total_loss: 1.061 loss_cls: 0.2716 loss_box_reg: 0.312 loss_rpn_cls: 0.1757 loss_rpn_loc: 0.2165 time: 2.1255 last_time: 2.1453 data_time:
[11/11 22:07:44 d2.utils.events]: eta: 0:39:47 iter: 319 total_loss: 1.099 loss_cls: 0.2666 loss_box_reg: 0.3316 loss_rpn_cls: 0.1616 loss_rpn_loc: 0.258 time: 2.1150 last_time: 2.1518 data_time:
[11/11 22:08:26 d2.utils.events]: eta: 0:39:14 iter: 339 total_loss: 0.7895 loss_cls: 0.1759 loss_box_reg: 0.2521 loss_rpn_cls: 0.1537 loss_rpn_loc: 0.213 time: 2.1156 last_time: 2.5653 data_time:
[11/11 22:09:14 d2.utils.events]: eta: 0:38:49 iter: 359 total_loss: 0.9343 loss_cls: 0.2964 loss_box_reg: 0.369 loss_rpn_cls: 0.1259 loss_rpn_loc: 0.1829 time: 2.1319 last_time: 2.6413 data_time:
[11/11 22:09:56 d2.utils.events]: eta: 0:38:12 iter: 379 total_loss: 0.8501 loss_cls: 0.2152 loss_box_reg: 0.2273 loss_rpn_cls: 0.145 loss_rpn_loc: 0.2228 time: 2.1205 last_time: 1.7900 data_time:
[11/11 22:10:37 d2.utils.events]: eta: 0:37:28 iter: 399 total_loss: 0.9636 loss_cls: 0.2667 loss_box_reg: 0.3844 loss_rpn_cls: 0.126 loss_rpn_loc: 0.2001 time: 2.1254 last_time: 1.5370 data_time:
[11/11 22:11:18 d2.utils.events]: eta: 0:36:44 iter: 419 total_loss: 0.9925 loss_cls: 0.2428 loss_box_reg: 0.3107 loss_rpn_cls: 0.1323 loss_rpn_loc: 0.1827 time: 2.1222 last_time: 1.9124 data_time:
[11/11 22:12:06 d2.utils.events]: eta: 0:36:10 iter: 439 total_loss: 1.02 loss_cls: 0.2044 loss_box_reg: 0.2899 loss_rpn_cls: 0.1372 loss_rpn_loc: 0.2331 time: 2.1335 last_time: 2.2892 data_time:
[11/11 22:12:50 d2.utils.events]: eta: 0:35:25 iter: 459 total_loss: 0.8739 loss_cls: 0.2137 loss_box_reg: 0.3273 loss_rpn_cls: 0.1204 loss_rpn_loc: 0.1864 time: 2.1358 last_time: 1.7495 data_time:
[11/11 22:13:36 d2.utils.events]: eta: 0:34:48 iter: 479 total_loss: 0.9252 loss_cls: 0.2576 loss_box_reg: 0.2901 loss_rpn_cls: 0.1102 loss_rpn_loc: 0.2052 time: 2.1431 last_time: 2.5612 data_time:
[11/11 22:14:19 d2.utils.events]: eta: 0:34:07 iter: 499 total_loss: 0.9243 loss_cls: 0.2151 loss_box_reg: 0.324 loss_rpn_cls: 0.1311 loss_rpn_loc: 0.197 time: 2.1431 last_time: 1.7906 data_time
```

- Similarly for 700, it converges to 0.9.


Low number of iterations: 1000

It was still converging, and the results were not promising with 1000 iterations. Loss was still reducing and could be further reduced with more iterations.

Visualization with 1000 iterations



Visualization with 1500 iterations



Part 2

Hyperparameter setting

- Number of epochs: 70
- Batch size: 8
- Learning rate: 0.01

Loss of first few iterations:

Epoch: 0, Loss: 0.4026719927787781
Epoch: 1, Loss: 0.2648690938949585
Epoch: 2, Loss: 0.23141880333423615
Epoch: 3, Loss: 0.21204398572444916
Epoch: 4, Loss: 0.19880929589271545
Epoch: 5, Loss: 0.18895624577999115

Loss of last few iterations:

Epoch: 64, Loss: 0.07312179356813431
Epoch: 65, Loss: 0.07243578881025314
Epoch: 66, Loss: 0.07232359796762466
Epoch: 67, Loss: 0.07169823348522186
Epoch: 68, Loss: 0.0712670236825943
Epoch: 69, Loss: 0.07036665081977844

Final Architecture:

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 8, 128, 128]	224
BatchNorm2d-2	[-1, 8, 128, 128]	16
ReLU-3	[-1, 8, 128, 128]	0
conv-4	[-1, 8, 128, 128]	0
Conv2d-5	[-1, 16, 128, 128]	1,168
BatchNorm2d-6	[-1, 16, 128, 128]	32
ReLU-7	[-1, 16, 128, 128]	0
conv-8	[-1, 16, 128, 128]	0
MaxPool2d-9	[-1, 16, 64, 64]	0
down-10	[-1, 16, 64, 64]	0
Conv2d-11	[-1, 32, 64, 64]	4,640
BatchNorm2d-12	[-1, 32, 64, 64]	64
ReLU-13	[-1, 32, 64, 64]	0
conv-14	[-1, 32, 64, 64]	0

MaxPool2d-15	[-1, 32, 32, 32]	0
down-16	[-1, 32, 32, 32]	0
Conv2d-17	[-1, 64, 32, 32]	18,496
BatchNorm2d-18	[-1, 64, 32, 32]	128
ReLU-19	[-1, 64, 32, 32]	0
conv-20	[-1, 64, 32, 32]	0
MaxPool2d-21	[-1, 64, 16, 16]	0
down-22	[-1, 64, 16, 16]	0
Conv2d-23	[-1, 128, 16, 16]	73,856
BatchNorm2d-24	[-1, 128, 16, 16]	256
ReLU-25	[-1, 128, 16, 16]	0
conv-26	[-1, 128, 16, 16]	0
MaxPool2d-27	[-1, 128, 8, 8]	0
down-28	[-1, 128, 8, 8]	0
Conv2d-29	[-1, 256, 8, 8]	295,168
BatchNorm2d-30	[-1, 256, 8, 8]	512
ReLU-31	[-1, 256, 8, 8]	0
conv-32	[-1, 256, 8, 8]	0
MaxPool2d-33	[-1, 256, 4, 4]	0
down-34	[-1, 256, 4, 4]	0
Conv2d-35	[-1, 512, 4, 4]	1,180,160
BatchNorm2d-36	[-1, 512, 4, 4]	1,024
ReLU-37	[-1, 512, 4, 4]	0
conv-38	[-1, 512, 4, 4]	0
MaxPool2d-39	[-1, 512, 2, 2]	0
down-40	[-1, 512, 2, 2]	0
ConvTranspose2d-41	[-1, 512, 4, 4]	1,049,088
Conv2d-42	[-1, 256, 4, 4]	1,179,904
BatchNorm2d-43	[-1, 256, 4, 4]	512
ReLU-44	[-1, 256, 4, 4]	0
conv-45	[-1, 256, 4, 4]	0
up-46	[-1, 256, 4, 4]	0
ConvTranspose2d-47	[-1, 256, 8, 8]	262,400
Conv2d-48	[-1, 128, 8, 8]	295,040
BatchNorm2d-49	[-1, 128, 8, 8]	256
ReLU-50	[-1, 128, 8, 8]	0
conv-51	[-1, 128, 8, 8]	0
up-52	[-1, 128, 8, 8]	0
ConvTranspose2d-53	[-1, 128, 16, 16]	65,664
Conv2d-54	[-1, 64, 16, 16]	73,792
BatchNorm2d-55	[-1, 64, 16, 16]	128
ReLU-56	[-1, 64, 16, 16]	0
conv-57	[-1, 64, 16, 16]	0
up-58	[-1, 64, 16, 16]	0
ConvTranspose2d-59	[-1, 64, 32, 32]	16,448
Conv2d-60	[-1, 32, 32, 32]	18,464
BatchNorm2d-61	[-1, 32, 32, 32]	64
ReLU-62	[-1, 32, 32, 32]	0
conv-63	[-1, 32, 32, 32]	0
up-64	[-1, 32, 32, 32]	0
ConvTranspose2d-65	[-1, 32, 64, 64]	4,128
Conv2d-66	[-1, 16, 64, 64]	4,624
BatchNorm2d-67	[-1, 16, 64, 64]	32
ReLU-68	[-1, 16, 64, 64]	0
conv-69	[-1, 16, 64, 64]	0
up-70	[-1, 16, 64, 64]	0
ConvTranspose2d-71	[-1, 16, 128, 128]	1,040

Conv2d-72	[-1, 4, 128, 128]	580
BatchNorm2d-73	[-1, 4, 128, 128]	8
ReLU-74	[-1, 4, 128, 128]	0
conv-75	[-1, 4, 128, 128]	0
up-76	[-1, 4, 128, 128]	0
Conv2d-77	[-1, 1, 128, 128]	37
conv-78	[-1, 1, 128, 128]	0

Total params: 4,547,953
 Trainable params: 4,547,953
 Non-trainable params: 0


Input size (MB): 0.19
 Forward/backward pass size (MB): 33.25
 Params size (MB): 17.35
 Estimated Total Size (MB): 50.79

Final Mean IOU: 0.86


#images: 7980, Mean IoU: 0.9057676527897925

Test plane image and predicted masks:


Sample 1



Sample 2



Sample 3



Part3

Group Member: Jashanraj Singh Gosain (301435386) and Archit Verma (301401441)

Kaggle Name : Jashanraj Singh Gosain, Archit

Score: 0.38

Sample1


Image



Ground Truth



Prediction




Sample 2


Image



Ground Truth



Prediction



Sample 3


Image



Ground Truth



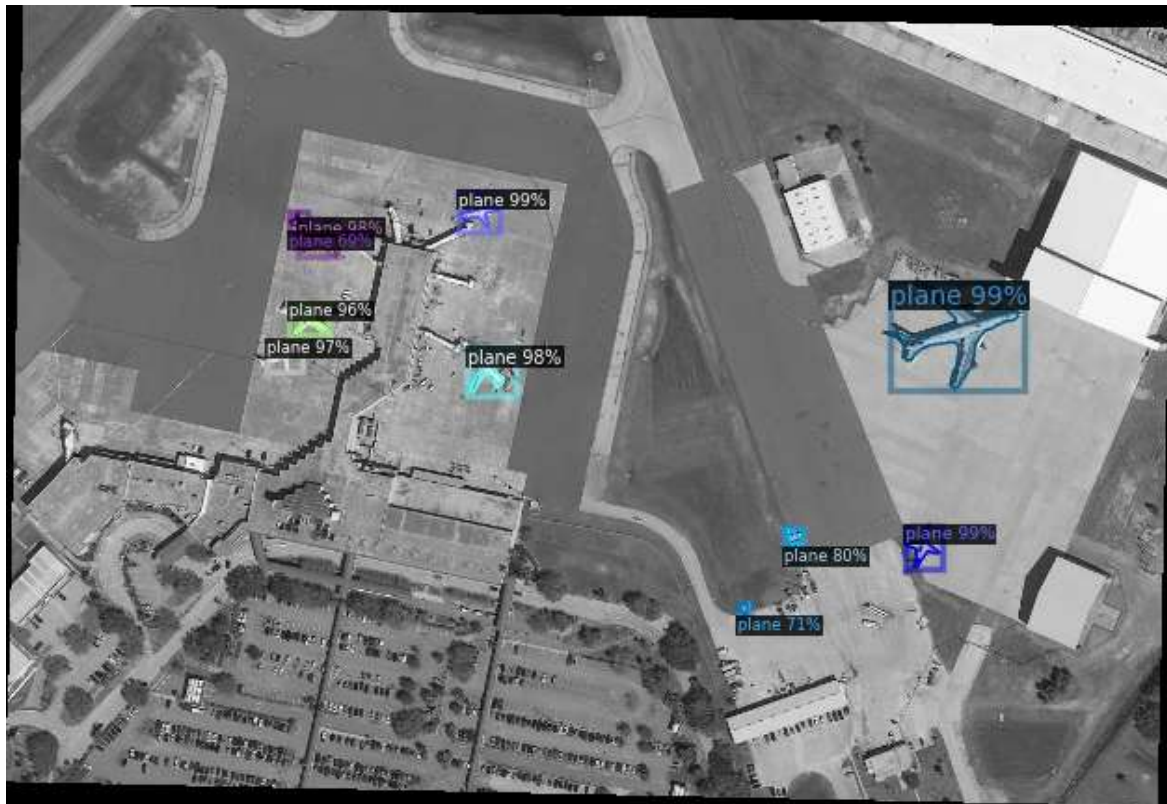
Prediction



Part 4

Visualization of test samples


Sample 1



Sample 2



Sample 3



Evaluation:

AP for bbox is 32 in part 4 which is lower than AP for bbox in part 2.

AP50 (51) ,AP75 (35), Aps (21), APm (41) and Apl (62) for Part 2 have higher value than it's corresponding values in part 4 (AP50 – 56, AP75 – 41, Aps – 26, APm – 45, Apl – 68) .

“mask_rcnn_R_50_FPN_3x.yaml”:

- It is computationally efficient as compared to “faster_rcnn_X_101_32x8d_FPN_3x.yaml”.
- It provides good results but with increased computational efficiency it's accuracy decreases.

“faster_rcnn_X_101_32x8d_FPN_3x.yaml”

- It is useful to calculate achieve high accuracy
- It's computational cost is higher than “mask_rcnn_R_50_FPN_3x.yaml”. so, it will need more time and resources to train.

```
[11/12 02:23:17 d2.evaluation.coco_evaluation]: Evaluation results for bbox:
| AP   | AP50 | AP75 | Aps  | APm  | Apl  |
|:-----|:-----|:-----|:-----|:-----|:-----|
| 32.011 | 51.126 | 35.685 | 21.095 | 41.038 | 62.045 |
Loading and preparing results...
DONE (t=0.04s)
creating index...
index created!
Running per image evaluation...
Evaluate annotation type *segm*
DONE (t=9.94s).
Accumulating evaluation results...
DONE (t=0.11s).
Average Precision (AP) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.091
Average Precision (AP) @[ IoU=0.50 | area= all | maxDets=100 ] = 0.300
Average Precision (AP) @[ IoU=0.75 | area= all | maxDets=100 ] = 0.025
Average Precision (AP) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.047
Average Precision (AP) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.101
Average Precision (AP) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.348
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 1 ] = 0.008
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets= 10 ] = 0.058
Average Recall (AR) @[ IoU=0.50:0.95 | area= all | maxDets=100 ] = 0.121
Average Recall (AR) @[ IoU=0.50:0.95 | area= small | maxDets=100 ] = 0.062
Average Recall (AR) @[ IoU=0.50:0.95 | area=medium | maxDets=100 ] = 0.145
Average Recall (AR) @[ IoU=0.50:0.95 | area= large | maxDets=100 ] = 0.401
[11/12 02:23:27 d2.evaluation.coco_evaluation]: Evaluation results for segm:
| AP   | AP50 | AP75 | Aps  | APm  | Apl  |
|:-----|:-----|:-----|:-----|:-----|:-----|
| 9.113 | 29.981 | 2.462 | 4.719 | 10.124 | 34.849 |
OrderedDict([('bbox', {'AP': 32.01118799298497, 'AP50': 51.12643736491969, 'AP75': 35.685389168440565,
```